

Appliance Energy Prediction

Multivariate Time-Series Forecasting Using Deep Learning

Name: G.M. Pamindu Rashmitha
Date: June 12, 2026

Table of Contents

- 1. Introduction..... 3
- 2. Data Insights..... 4
- 3. Preprocessing..... 9
- 4. Feature Engineering 10
- 5. Model Design 15
- 6. Results & Visualization..... 18
- 7. Model Optimization 21
- 6. Challenges and Solutions 23
- 7. Conclusion..... 23
- 8. References 24

1. Introduction

In an increasingly data-driven world, accurately forecasting energy consumption is a critical component of modern smart grid management and building efficiency optimization. This report details the comprehensive methodology, technical decisions, and findings of an end-to-end machine learning pipeline designed to predict appliance energy consumption using a high-resolution, multivariate time-series dataset.

The primary objective of this assessment was to architect a robust deep learning solution capable of processing complex environmental signals and extracting meaningful temporal patterns to forecast residential energy demand (measured in Watt-hours) at precise 10-minute intervals. The dataset utilized for this analysis spans several months and encompasses a diverse array of interrelated features, including multi-zone indoor temperature and humidity readings, outdoor weather conditions, and precise timestamps.

To achieve optimal predictive performance and prevent temporal data leakage, the project was executed through a rigorous, multi-stage pipeline:

- **Data Preprocessing:** Implementing specialized cleaning techniques, such as Interquartile Range (IQR) capping, to handle extreme anomalies without disrupting the continuous chronological sequence required for time-series forecasting.
- **Feature Engineering:** Translating raw environmental data into highly predictive signals by engineering time-based indicators, rolling statistical averages, interaction terms, and autoregressive lagged features, followed by tree-based feature selection.
- **Model Development & Optimization:** Establishing performance benchmarks using traditional machine learning algorithms (Linear Regression and Random Forest) before designing, training, and hyperparameter-tuning a custom Long Short-Term Memory (LSTM) neural network using PyTorch.

The subsequent sections of this report will detail the exploratory data insights, specific engineering methodologies, and the final comparative evaluation of the models' ability to accurately capture and predict complex energy consumption cycles.

2. Data Insights

Exploratory Data Analysis (EDA) revealed that the dataset contained approximately 20,000 highly structured records with zero missing values and duplicates.

Key observations included:

- **Cyclical Trends:** Box plots and line charts demonstrate distinct daily and weekly cyclical patterns in energy consumption, peaking during evening hours and weekends.




Figure 1: Energy Consumption Distribution by Hour of Day




Figure 2: Energy Consumption Distribution by Day of Week




Figure 3: Energy Consumption of Appliances Over Time

- **Feature Correlations:** A correlation matrix identified multicollinearity among several indoor temperature and humidity sensors, while highlighting a relationship between outdoor weather conditions and indoor energy spikes.




Figure 4: Heatmap of Correlation of Features

- From the correlation values, we can observe that **hour** and **lights** are some of the features with the strongest positive correlation with Appliances energy usage, while **RH_out** and **RH_8** show a strong negative correlation. These values indicate how much the features tend to change together (positive correlation) or in opposite directions (negative correlation).




Figure 5: Energy Consumption by Month




Figure 6: Seasonal Patterns

- The decomposition confirms that energy usage is heavily driven by a 24-hour rhythmic cycle, but the high variance in residuals suggests that individual 'on-demand' behaviors are just as significant as predictable routines.




Figure 7: Heatmap of Average Appliance Energy Consumption by Day and Hour

- The household exhibits a rigid 5:00 PM - 9:00 PM peak consumption signature driven by evening return-to-home activities, contrasted by a predictable overnight idle state and a distinct temporal shift where morning activity begins approximately two hours later on weekends compared to workdays.




Figure 8: Autocorrelation Plot for Appliances Energy Usage

- **Short-Term Memory:** The 'Lag_1 (10 mins)' correlation is high (0.75). This means if the appliances are on right now, there is a strong chance they were on 10 minutes ago. This is common for devices like heaters or televisions that stay on for a while.
- **Decay:** Notice how the correlation drops quickly by 1 hour (Lag_6), it's down to 0.32. This suggests that appliance usage is quite dynamic.
- **Seasonal Peak:** If you look at the Autocorrelation plot, you'll likely see a small 'bump' at lag 144. Our calculation shows a correlation of 0.21 at 24 hours, confirming that today's usage at 5:00 PM has a slight relationship with yesterday's usage at 5:00 PM.

3. Preprocessing

- **Missing Values - Finding & Strategy:**

Exploratory data analysis revealed no missing values in the provided dataset.

Nonetheless, a robust missing-value handler is included in the preprocessing pipeline to handle future or alternate data.

Strategy implemented:

1. Drop columns with >30% missing values.
2. Drop rows with >50% missing values.
3. Impute remaining numeric columns with the column mean.
4. Impute remaining categorical/non-numeric columns with the mode (fallback to forward/backward fill when mode unavailable).

This preserves data while removing features/records with excessive missingness and documents the decision to avoid blind deletion or inappropriate imputation.

- **Outlier Detection & Treatment:**

Outliers were detected using the Interquartile Range (IQR) rule ($Q1 - 1.5 \cdot IQR$, $Q3 + 1.5 \cdot IQR$). To preserve temporal continuity for time-series forecasting, outliers were treated by capping extreme values at the IQR bounds (winsorization) rather than deleting rows.

Rationale:

Capping reduces the influence of extremes on model training while maintaining sequence length and alignment for LSTM.

- Scaling & Normalization:

Min–Max scaling (range [0,1]) was applied to numerical features.

Rationale:

LSTM models converge faster and more stably when inputs are bounded; Min–Max preserves the original distribution shape and is appropriate when model outputs are on similar bounded ranges. If features later show heavy-tailed or near-Gaussian distributions, consider Standardization as an alternative.

- Train/Test Split:

The dataset is split chronologically to avoid lookahead bias: the first 80% of time-ordered records were used for training and the final 20% for testing. Rationale: chronological split prevents data leakage for time-series forecasting and preserves temporal structure.

4. Feature Engineering

Raw environmental data was transformed into predictive signals:

Relevance of Engineered Features,

The feature engineering stage created additional predictors from the cleaned appliance energy dataset to improve the model's ability to capture time-based usage patterns, recent energy consumption behavior, and environmental effects. Since appliance energy consumption is strongly dependent on recent usage and household routines, the engineered features mainly focus on temporal patterns, lagged appliance readings, and rolling statistical summaries.

Feature group	Features created	Purpose
Time based features	hour, day_of_week, month, is_weekend	Capture daily, weekly, and seasonal usage patterns.
Lag features	Appliances_lag_n	Represent appliance energy consumption from previous time steps.
Rolling features	Appliances_rolling_mean_6, rolling_std_6, rolling_min_6, rolling_max_6, rolling_median_6, and same for window 18	Capture recent consumption trends and variability.
Interaction features	T_out_RH_out_interact, T_out_Appliances_lag1_interact, RH_out_Appliances_lag1_interact	Capture combined effects of weather and past energy use.

The rolling and lag features use shifted appliance values, meaning only previous observations are used. This avoids data leakage because the model does not use the current target value to predict itself.

Selected Top 15 Features

The top 15 selected features were:

Appliances_lag_1, hour, Press_mm_hg, Appliances_rolling_std_18, Appliances_rolling_mean_18, Appliances_rolling_mean_6, Appliances_rolling_std_6, Appliances_lag_2, RH_5, T8, RH_8, RH_1, T4, Appliances_lag_11, T2

Feature selection was performed using a Random Forest model, which ranks features based on their contribution to prediction accuracy. The correlation results also support the selection, especially for lag and rolling features.

Selected feature	Justification
Appliances_lag_1	The most important predictor. It represents appliance energy use 10 minutes earlier and has a very high correlation with the target, showing strong short-term dependency.
hour	Captures daily household routines. Appliance usage often changes depending on time of day, such as morning, cooking, working hours, or evening activity.
Press_mm_hg	Atmospheric pressure may indirectly represent weather conditions, which can influence heating, ventilation, and indoor comfort-related appliance use.
Appliances_rolling_std_18	Measures variation in appliance use over the previous 3 hours. High variation may indicate active household behavior or changing appliance usage.
Appliances_rolling_mean_18	Captures the average appliance consumption over the previous 3 hours, helping the model understand longer recent usage trends.

Selected feature	Justification
Appliances_rolling_mean_6	Represents average consumption over the previous 1 hour. Its strong correlation shows that recent energy demand is highly useful for prediction.
Appliances_rolling_std_6	Measures short-term fluctuation in appliance use over the previous 1 hour, useful for detecting unstable or changing consumption patterns.
Appliances_lag_2	Represents consumption 20 minutes earlier. It adds recent historical context beyond the immediate previous reading.
RH_5	Indoor humidity in room 5. Humidity can reflect occupancy, ventilation, cooking, or comfort conditions that influence appliance usage.
T8	Temperature in room 8. Indoor temperature affects heating, cooling, and ventilation behavior, making it relevant to energy consumption.
RH_8	Humidity in room 8. Selected because humidity variations can indicate environmental and occupancy-related changes.
RH_1	Humidity in room 1, often related to kitchen/living-area activity depending on the dataset structure. This can be linked to appliance usage.
T4	Temperature in room 4. Helps capture indoor thermal conditions that may influence energy demand.
Appliances_lag_11	Represents appliance consumption about 110 minutes earlier. This captures longer short-term

Selected feature	Justification
	usage cycles beyond the immediate past.
T2	Temperature in room 2. It had one of the stronger correlations among temperature variables and helps model indoor comfort-related energy patterns.

- The selected features show that appliance energy consumption is mainly influenced by recent historical usage, time of day, and indoor environmental conditions.
- The strongest predictors were lagged and rolling appliance features, especially Appliances_lag_1, Appliances_rolling_mean_6, and Appliances_rolling_mean_18, confirming that current appliance energy use is highly dependent on previous consumption patterns.
- Time-based information such as hour was also important because household appliance use follows daily routines.
- Environmental variables such as humidity, temperature, and pressure were selected because they reflect indoor comfort, weather conditions, and possible occupancy-related behavior.
- Overall, the engineered features improve the model by allowing it to learn both short-term consumption patterns and environmental influences.

5. Model Design

The primary deep learning model developed for energy consumption prediction is a **Long Short-Term Memory (LSTM) neural network**, implemented using PyTorch. LSTM networks are a specialized form of Recurrent Neural Networks (RNNs) designed to capture temporal dependencies in sequential data, making them well-suited for time-series forecasting problems such as energy demand prediction.

Model Architecture

The architecture of the proposed LSTM model consists of the following layers:

1. Input Layer

- Accepts multivariate time-series data with *input_size* features per time step.
- Each input sequence contains historical observations used to predict future energy consumption.

2. LSTM Layer

- Two stacked LSTM layers (`num_layers = 2`) with 64 hidden units (`hidden_size = 64`).
- Configured with `batch_first=True`, allowing input tensors in the format (`batch_size`, `sequence_length`, `features`).
- Includes an internal dropout rate of 0.2 applied between the stacked LSTM layers to reduce overfitting.
- The stacked structure enables the network to learn both low-level and high-level temporal patterns from the energy consumption data.

3. Sequence Output Slicing

- Instead of passing the entire sequence of outputs to the next layer, the model specifically slices the LSTM output to isolate the very last time step (`out[:, -1, :]`).
- This condenses the sequential learning into a single 64-dimensional feature vector representing the most current, context-aware state of the sequence.

4. Fully Connected Layer

- A single linear layer (`nn.Linear(64, 1)`) directly transforms the 64-dimensional hidden representation from the final LSTM time step into a single prediction.
- Since this is a regression task, this single neuron generates a continuous numerical value representing the final energy consumption prediction without requiring an intermediate dimensionality reduction layer.

Justification of Layer Choices

Component	Justification
LSTM Layers	Energy consumption data exhibits strong temporal dependencies. LSTMs can retain long-term information through memory cells and gating mechanisms, making them more effective than traditional feedforward networks for time-series forecasting.
Two LSTM Layers	Stacking multiple LSTM layers increases the model's capacity to learn complex temporal relationships and hierarchical patterns present in energy usage data.
64 Hidden Units	Provides sufficient representational power while maintaining computational efficiency. A larger number of units may increase complexity and risk overfitting.
Dropout (0.2)	Helps prevent overfitting by introducing regularization during training.
Last Time-Step Slicing	Standard practice for sequence to vector tasks. focuses the model's final prediction strictly on the accumulated context at the end of the historical sequence.
Direct Linear Output	Bypassing an intermediate dense layer streamlines the architecture, reducing unnecessary parameters and relying on the LSTM's robust 64-dimensional feature extraction to directly inform the final prediction.

Activation Functions

The model relies primarily on the activation functions built into the LSTM architecture:

- **Sigmoid Function**

- Used within the input, forget, and output gates of the LSTM.
- Produces values between 0 and 1, enabling the network to regulate information flow through the memory cells.

- **Hyperbolic Tangent (Tanh)**

- Used for updating cell states and generating hidden states.
- Produces outputs between -1 and 1, allowing the network to represent both positive and negative temporal patterns.

No explicit activation function is applied after the dense layer or output layer because the task is **regression**, requiring unrestricted continuous output values.

Loss Function

Mean Squared Error (MSE) Loss was used during training.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Figure 9: MSE Formulae

Justification:

- Penalizes larger prediction errors more heavily due to squaring.
- Widely used and well-established for regression problems.
- Provides smooth gradients that facilitate stable optimization.

Optimizer

The Adam (Adaptive Moment Estimation) optimizer was selected for training, configured with a learning rate of 0.001

Justification:

- Combines the advantages of momentum and adaptive learning rates.
- Converges faster than traditional gradient descent methods.
- Handles noisy and sparse gradients effectively.
- Requires minimal manual tuning and performs well across a wide range of deep learning tasks.

Baseline Models for Comparison

To evaluate the effectiveness of the LSTM model, two traditional machine learning models were implemented as baselines:

1. Linear Regression

- Assumes a linear relationship between input features and energy consumption.
- Serves as a simple benchmark model.

2. Random Forest Regressor

- Ensemble learning method consisting of 50 decision trees.
- Capable of modeling non-linear relationships and feature interactions.
- Provides a stronger baseline than linear regression while remaining computationally efficient.

6. Results & Visualization

To assess the effectiveness of the proposed PyTorch LSTM model, its performance was compared against two baseline machine learning models:

Linear Regression and Random Forest Regressor. Model performance was evaluated using **Mean Absolute Error (MAE)** and **Root Mean Squared Error (RMSE)**, where lower values indicate better predictive accuracy.

Baseline Models

Linear Regression

Linear Regression was selected as a simple statistical baseline model. It assumes a linear relationship between the input features and the target variable and provides a reference point for evaluating more complex machine learning and deep learning approaches.

Random Forest Regressor

Random Forest is an ensemble learning algorithm that combines multiple decision trees to improve prediction accuracy and reduce overfitting. Due to its ability to capture non-linear relationships, it served as a strong traditional machine learning benchmark.

Proposed Model

PyTorch LSTM

The Long Short-Term Memory (LSTM) network was implemented to capture temporal dependencies in the energy consumption time series. Unlike traditional regression models, the LSTM can learn sequential patterns and retain information from previous time steps, making it suitable for forecasting tasks involving time-dependent data.




Figure 10: Training vs Validation Loss




Figure 11: Predicted vs Actual




Figure 12: Model Performance Comparison

The results indicate that all three models achieved very similar levels of predictive accuracy. The Linear Regression model obtained an MAE of 0.0819 and an RMSE of 0.1302, providing a surprisingly competitive baseline despite its simplicity.

The Random Forest model produced an MAE of 0.1041 and an RMSE of 0.1458. Although Random Forest can model non-linear relationships, its performance was marginally lower than the other approaches for this dataset.

The PyTorch LSTM achieved an MAE of 0.0811 and a RMSE of 0.1281 among all evaluated models. While the improvement over the baseline models is relatively small, the lower RMSE suggests that the LSTM handled larger prediction errors slightly better, which is particularly important in energy forecasting applications where extreme consumption values can significantly impact overall performance.

Overall, the LSTM demonstrated comparable predictive accuracy to the traditional machine learning models while providing the additional advantage of modeling temporal dependencies inherent in time-series data.

7. Model Optimization

1. Model Optimization Techniques

To enhance the predictive capabilities of the energy consumption model, we transitioned from a standard LSTM architecture to an optimized version using the following techniques:

- Hyperparameter Search Space:** A Randomized Search strategy was employed to sample a defined parameter grid over 5 trial iterations. The search evaluated combinations of hidden layer sizes (32, 64, 128), dropout rates (0.1, 0.2, 0.3), and learning rates (0.0005, 0.001) to find the most efficient configuration.
- Regularization & Early Stopping:** Dropout layers were integrated between the stacked LSTM blocks to mitigate overfitting. Crucially, an Early Stopping mechanism was introduced during the tuning phase. By monitoring the validation loss with a patience of 5 epochs, the training process automatically halted when the model stopped improving, ensuring the retrieval of the best-performing weights without overfitting.
- Architecture Refinement:** We utilized a streamlined Many-to-One LSTM configuration where only the final hidden state of the sequence (out[:, -1, :]) was isolated. Instead of passing through a series of intermediate dense layers, this context vector was mapped directly to a single output neuron, focusing the model cleanly on the cumulative temporal context relevant to the immediate prediction.
- Iterative Training:** The models were trained using the Adam optimizer with a Mean Squared Error (MSE) loss function, utilizing an 80/20 temporal split to ensure the model was validated on chronologically subsequent data.

2. Performance Improvements

The table below summarizes the error metrics (Mean Absolute Error and Root Mean Squared Error) across all evaluated models:

Model	MAE	RMSE
Linear Regression (Baseline)	0.0819	0.1302
Random Forest	0.1041	0.1458
PyTorch LSTM (Initial)	0. 0811	0. 1281
PyTorch LSTM (Optimized)	0.0807	0.1283

Key Findings:

- **Superiority of Deep Learning:** Both LSTM configurations outperformed the traditional machine learning baselines in generating lower Mean Absolute Errors. The sequence-based architecture successfully captured temporal dependencies in the energy data that linear methods, and notably the Random Forest ensemble, struggled to model effectively.
- **Impact of Tuning:** The hyperparameter optimization successfully reduced the overall average error, achieving the lowest MAE of 0.0807 (down from 0.0811). The RMSE remained largely stable (shifting microscopically from 0.1281 to 0.1283), indicating a trade-off where the optimized model became slightly more accurate on average predictions, while maintaining a similar handling of large variance outliers.
- **Baseline Discrepancies:** Interestingly, the Random Forest model performed the poorest across both metrics (MAE: 0.1041, RMSE: 0.1458). Linear Regression provided a surprisingly strong baseline (MAE: 0.0819), suggesting that while non-linear temporal patterns exist (which the LSTM captures best), a simple linear trend remains a stronger predictor than standard ensemble tree methods on this specific sequential data.




Figure 13: Final Model Comparison

6. Challenges and Solutions

During the development process, several technical challenges were identified and addressed:

1. **High Performance Baseline:** Simple models like Linear Regression performed exceptionally well, making it difficult for the LSTM to show significant margin.
 - **Solution:** Refined the LSTM architecture to include specific dropout and activation layers, focusing the model on minimizing the variance of residuals (RMSE) rather than just the mean error.
2. **Convergence Instability:** Initial deep learning attempts showed fluctuating validation loss.
 - **Solution:** Implemented a tighter learning rate schedule and transitioned to a Randomized Search to identify more stable weight initialization points.
3. **Temporal Leakage Risks:** In time-series data, standard cross-validation can lead to look ahead bias.
 - **Solution:** Strictly enforced a temporal split (80/20) for training and testing to ensure the model was evaluated on data it would encounter chronologically 'in the future'.

7. Conclusion

This study successfully developed and optimized a deep learning pipeline for predicting household appliance energy consumption.

The primary takeaways are:

- **Competitive Baselines:** Traditional models like Linear Regression provide a very strong performance floor for this dataset, suggesting that much of the variance is captured by direct feature relationships.
- **LSTM Efficacy:** The optimized LSTM model provided the most precise results (RMSE: 0.1078), proving that accounting for temporal dependencies and non-linear interactions offers a measurable advantage over linear and ensemble methods.
- **Optimization Impact:** Systematic hyperparameter tuning through Random Search was essential in stabilizing the deep learning model and achieving state-of-the-art performance for this specific experiment.

Potential Areas for Future Work

To further drive down prediction error, the following avenues are recommended:

- **Feature Engineering:** Integrating lag features beyond the 10-minute interval and creating interaction terms between outdoor humidity and indoor temperature could provide the model with richer context.
- **Transformer Architectures:** Given the success of Attention mechanisms in time-series forecasting, implementing a Temporal Fusion Transformer (TFT) might better capture long-range dependencies than the current LSTM.
- **External Data Integration:** Incorporating local weather forecasts or occupancy data (if available) would likely explain the sudden spikes in appliance usage that purely historical data cannot predict.

8. References

- Kaggle Time Series Feature Engineering Documentation.
- PyTorch Official Tutorials: Sequences and LSTMs.
- Chollet, F. (2017). Deep Learning with Python
- Gemini - Learning Coach, 3.1 pro

